

```
In [83]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

1. LINEAR REGRESSION ON TIME SERIES DATA

```
In [84]: df = pd.read_csv('time_series.csv')
df.head()
```

```
Out[84]:
```

	date	X1	X2	X3	X4	X5	X6	Y
0	2016-07-01 00:00:00	5.827	2.009	1.599	0.462	4.203	1.340	30.531000
1	2016-07-01 01:00:00	5.693	2.076	1.492	0.426	4.142	1.371	27.787001
2	2016-07-01 02:00:00	5.157	1.741	1.279	0.355	3.777	1.218	27.787001
3	2016-07-01 03:00:00	5.090	1.942	1.279	0.391	3.807	1.279	25.044001
4	2016-07-01 04:00:00	5.358	1.942	1.492	0.462	3.868	1.279	21.948000

```
In [85]: sample_df = df.head(300)
```

```
In [86]: fig, axes = plt.subplots(6, 1, figsize=(12, 10), sharex=True)
```

```
# Loop to plot each variable
lst = ['X1', 'X2', 'X3', 'X4', 'X5', 'X6']
for i in range(0,6):
    sns.lineplot(x=lst[i], y='Y', data=df, ax=axes[i])
    axes[i].set_title(f'Time Series of {lst[i]}')
    axes[i].set_ylabel(var)

plt.tight_layout(rect=[0, 0.03, 1, 0.95])
```

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):


C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):



Exploring the Correlation

```
In [87]: corr_df = df.drop('date',axis=1)
corr_df.corr()
```

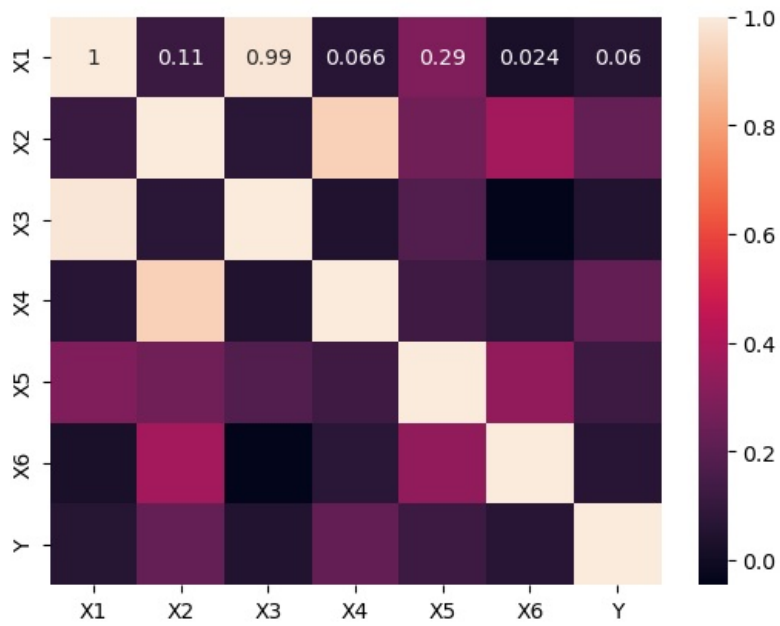
```
Out[87]:
```

	X1	X2	X3	X4	X5	X6	Y
X1	1.000000	0.114672	0.987355	0.066002	0.291418	0.023606	0.059916
X2	0.114672	1.000000	0.068817	0.930491	0.259487	0.377641	0.224354
X3	0.987355	0.068817	1.000000	0.046266	0.177491	-0.046519	0.050854
X4	0.066002	0.930491	0.046266	1.000000	0.128607	0.069419	0.220004
X5	0.291418	0.259487	0.177491	0.128607	1.000000	0.334563	0.118836
X6	0.023606	0.377641	-0.046519	0.069419	0.334563	1.000000	0.067455
Y	0.059916	0.224354	0.050854	0.220004	0.118836	0.067455	1.000000

Columns X1 and X3 look to be very closely correlated


```
In [88]: sns.heatmap(corr_df.corr(),annot=True)
```

```
Out[88]: <Axes: >
```



Plotting the Kernel Density

```
In [89]: df.plot(kind='kde')
plt.xlabel(" X Values")
plt.ylabel("Density")
plt.grid(True)
plt.show()
```



In []:

b.) Train/Test Split

```
In [90]: df
```

Out[90]:

	date	X1	X2	X3	X4	X5	X6	Y
0	2016-07-01 00:00:00	5.827	2.009	1.599	0.462	4.203	1.340	30.531000
1	2016-07-01 01:00:00	5.693	2.076	1.492	0.426	4.142	1.371	27.787001
2	2016-07-01 02:00:00	5.157	1.741	1.279	0.355	3.777	1.218	27.787001
3	2016-07-01 03:00:00	5.090	1.942	1.279	0.391	3.807	1.279	25.044001
4	2016-07-01 04:00:00	5.358	1.942	1.492	0.462	3.868	1.279	21.948000
...
17415	2018-06-26 15:00:00	-1.674	3.550	-5.615	2.132	3.472	1.523	10.904000
17416	2018-06-26 16:00:00	-5.492	4.287	-9.132	2.274	3.533	1.675	11.044000
17417	2018-06-26 17:00:00	2.813	3.818	-0.817	2.097	3.716	1.523	10.271000
17418	2018-06-26 18:00:00	9.243	3.818	5.472	2.097	3.655	1.432	9.778000
17419	2018-06-26 19:00:00	10.114	3.550	6.183	1.564	3.716	1.462	9.567000

17420 rows × 8 columns

In [91]:

```
train = df.loc[df['date'] <= '2017-06-26 23:00:00']
val = df.loc[(df['date'] > '2017-06-26 23:00:00') & (df['date'] <= '2017-10-24 23:00:00')]
test = df.loc[(df['date'] > '2017-10-24 23:00:00') & (df['date'] <= '2018-02-21 23:00:00')]

print("Training set size:", len(train))
print("Validation set size:", len(val))
print("Test set size:", len(test))
```

Training set size: 8664
Validation set size: 2879
Test set size: 2880

In [92]:

```
df['date'] = pd.to_datetime(df['date'])


train_end_date = '2017-06-26'
val_end_date = '2017-10-24'
test_end_date = '2018-02-21'

plt.figure(figsize=(12, 6))
plt.plot(df['date'], df['Y'], label='Target Variable (Y)')

plt.axvline(pd.to_datetime(train_end_date), color='green', linestyle='--', label='Train End')
plt.axvline(pd.to_datetime(val_end_date), color='orange', linestyle='--', label='Validation End')
plt.axvline(pd.to_datetime(test_end_date), color='red', linestyle='--', label='Test End')

plt.title("Time Series with Train, Validation, and Test Splits")
plt.xlabel("Date")
plt.ylabel("Target Variable (Y)")

plt.tight_layout()
plt.show()
```



In []:

In []:

c. Scaling

```
In [93]: def scale_variables(data, metrics):
        scaled_data = data.copy()

        for variable in metrics:
            mean = train[variable].mean()
            std = train[variable].std()
            scaled_data[variable] = (data[variable] - mean) / std

        return scaled_data

variables = ['X1', 'X2', 'X3', 'X4', 'X5', 'X6']

scaled_train = scale_variables(train, variables)
scaled_val = scale_variables(val, variables)
scaled_test = scale_variables(test, variables)
```

In []:

In []:

d. LSE

```
In [131]: X = scaled_train[['X1', 'X2', 'X3', 'X4', 'X5']]
        Y = scaled_train['Y']

        X = np.array(X)
        Y1 = np.array(Y)

        X_b = np.c_[np.ones(X.shape[0]), X] ### This will add an intercept column
```

```
In [133]: beta_hat = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(Y1)
        print("The estimated Parameters for Train are :", beta_hat)

        Y_pred1 = X_b.dot(beta_hat)
        print("Predictions for Train are :", Y_pred1)

        # Calculate Mean Squared Error
        mse = np.mean((Y1 - Y_pred1) ** 2)
        mae = np.mean(np.abs(Y1 - Y_pred1))
        print("Mean Squared Error and Mean Absolute Error for Train are :", mse, mae)
```

The estimated Parameters for Train are : [17.13662789 -1.91639749 5.26124611 1.78251227 -0.13756652 2.01455169]
Predictions for Train are : [19.49669908 19.55754778 18.1072748 ... 14.8454925 16.28669993 17.33686072]
Mean Squared Error and Mean Absolute Error for Train are : 50.894201122160645 5.476539015669759

In []:

Now applying to Validation and Test Dataset

```
In [134]: X = scaled_val[['X1', 'X2', 'X3', 'X4', 'X5']]
        Y = scaled_val['Y']

        X = np.array(X)
        Y2 = np.array(Y)
        X_b = np.c_[np.ones(X.shape[0]), X]

        beta_hat = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(Y2)
        print("The estimated Parameters for Validation are :", beta_hat)

        Y_pred2 = X_b.dot(beta_hat)
        print("Predictions for Validation:", Y_pred2)

        # Calculate Mean Squared Error
        mse = np.mean((Y2 - Y_pred2) ** 2)
        mae = np.mean(np.abs(Y - Y_pred2))
        print("Mean Squared Error and Mean Absolute Error for Validation:", mse, mae)
```

The estimated Parameters for Validation are : [13.81194261 7.68653594 -3.10614367 -7.71847754 2.83733242 1.98827336]
Predictions for Validation: [14.97764 14.20710219 14.49630545 ... 15.54576055 12.2865472 13.02299611]
Mean Squared Error and Mean Absolute Error for Validation: 11.195550354114475 2.7126178489513086

```
In [135]: X = scaled_test[['X1', 'X2', 'X3', 'X4', 'X5']]
```

```

Y = scaled_test['Y']

X = np.array(X)
Y3 = np.array(Y)
X_b = np.c_[np.ones(X.shape[0]), X]

beta_hat = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(Y3)
print("The estimated Parameters for Test are :", beta_hat)

Y_pred3 = X_b.dot(beta_hat)
print("Predictions for Test:", Y_pred3)

# Calculate Mean Squared Error
mse = np.mean((Y3 - Y_pred3) ** 2)
mae = np.mean(np.abs(Y - Y_pred3))
print("Mean Squared Error and Mean Absolute Error for Test:", mse,mae)

```

The estimated Parameters for Test are : [5.3615342 -4.56222556 3.66057755 3.88489226 -4.12782807 -0.67826609]
Predictions for Test: [5.86041621 4.66510387 6.64783407 ... 1.88372612 1.9110887 2.60269018]
Mean Squared Error and Mean Absolute Error for Test: 7.285752203254583 2.1457949999478316

```

In [142...] sns.lineplot(x=lst[i], y='Y', data=df)
#plt.scatter(X, exponential_func(X, *popt), label='Fitted Curve', color='red')


```

C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):
C:\Users\tegbe\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):

```

Out[142...] <Axes: xlabel='X6', ylabel='Y'>

```



```

In [140...] X_train = scaled_train[['X1', 'X2', 'X3', 'X4', 'X5']]
X_val = scaled_val[['X1', 'X2', 'X3', 'X4', 'X5']]
X_test = scaled_test[['X1', 'X2', 'X3', 'X4', 'X5']]

Y_train = scaled_train['Y']
Y_val = scaled_val['Y']
Y_test = scaled_test['Y']

Y_pred1
Y_pred2
Y_pred3

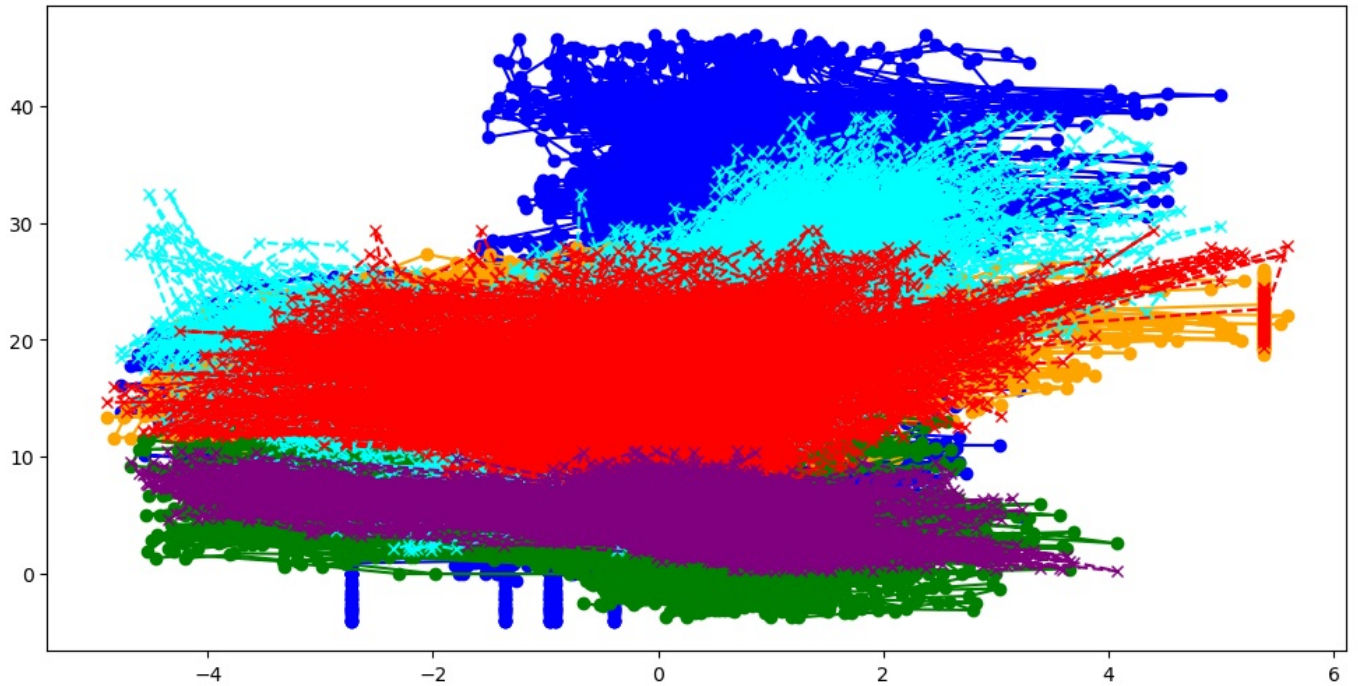
plt.figure(figsize=(12, 6))

# Plot true values
plt.plot(X_train, Y_train, marker='o', label='Train True', color='blue')
plt.plot(X_val, Y_val, marker='o', label='Validation True', color='orange')
plt.plot(X_test, Y_test, marker='o', label='Test True', color='green')

# Plot predicted values
plt.plot(X_train, Y_pred1, marker='x', linestyle='--', label='Train Predicted', color='cyan')
plt.plot(X_val, Y_pred2, marker='x', linestyle='--', label='Validation Predicted', color='red')
plt.plot(X_test, Y_pred3, marker='x', linestyle='--', label='Test Predicted', color='purple')

```

```
Out[148... [<matplotlib.lines.Line2D at 0x13b78fc7b90>,
<matplotlib.lines.Line2D at 0x13b78fc7bd0>,
<matplotlib.lines.Line2D at 0x13b755fb8d0>,
<matplotlib.lines.Line2D at 0x13b78fd0690>,
<matplotlib.lines.Line2D at 0x13b78fd0ad0>]
```



```
In [ ]:
```

```
In [ ]:
```

2. VISUALIZING LSE ON A TOY DATASET

```
In [143... df1 = pd.read_csv('toy_data.csv')
df1.head()
```

```
Out[143...
   X1    X2    Y
0  5.530492  8.136530  53.470131
1  5.111720  0.846906  15.925409
2  9.011047  6.510469  54.649639
3  7.806497  0.349096  24.003095
4  2.047190  1.057417  14.739897
```

```
In [149... X = df1[['X1', 'X2']]
Y = df1['Y']

#Converting it to Arrays
X = np.array(X)
Y = np.array(Y)

X_b = np.c_[np.ones(X.shape[0]), X] ### This will add an intercept column

### Now calculating the parameters using LSE

beta_hat = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(Y)
print("The estimated Parameters for Test are :", beta_hat)
```

The estimated Parameters for Test are : [3.66345247 1.96836519 4.93212246]

```
In [157... Y_pred = X_b.dot(beta_hat)
```

```
In [169... from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure(figsize=(10, 7))
ax = fig.add_subplot(111, projection='3d')


# Scatter plot
ax.scatter(df1['X1'], df1['X2'], df1['Y'], color='purple')
ax.scatter(df1['X1'], df1['X2'], Y_pred, color='red')

ax.set_xlabel('X1')
```



```
ax.set_ylabel('X2')  
ax.set_zlabel('Y')  
ax.set_title('3D Scatter Plot of X1 and X2')  
  
plt.show()
```

3D Scatter Plot of X1 and X2



In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js