

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from mpl_toolkits.mplot3d import Axes3D
from sklearn.metrics import confusion_matrix
```

```
In [422]: df = pd.read_csv('logistic.csv')
df.head()
```

```
Out[422]:
```

| | Y | X1 | X2 | X3 | X4 | X5 | X6 | X7 | X8 | X9 | ... | X21 | X22 | X23 | X24 | X25 | X26 | ... |
|---|---|-------|-------|--------|--------|---------|---------|--------|---------|--------|-----|-------|-------|--------|--------|--------|--------|-----|
| 0 | M | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14710 | 0.2419 | ... | 25.38 | 17.33 | 184.60 | 2019.0 | 0.1622 | 0.6656 | 0.7 |
| 1 | M | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07017 | 0.1812 | ... | 24.99 | 23.41 | 158.80 | 1956.0 | 0.1238 | 0.1866 | 0.2 |
| 2 | M | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.12790 | 0.2069 | ... | 23.57 | 25.53 | 152.50 | 1709.0 | 0.1444 | 0.4245 | 0.4 |
| 3 | M | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10520 | 0.2597 | ... | 14.91 | 26.50 | 98.87 | 567.7 | 0.2098 | 0.8663 | 0.6 |
| 4 | M | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10430 | 0.1809 | ... | 22.54 | 16.67 | 152.20 | 1575.0 | 0.1374 | 0.2050 | 0.4 |

5 rows × 31 columns

```
In [5]: ### Function to Split data into Training and Test set

def train_test_split(data):
    # Calculate the split index for 80% of the data
    split_idx = int(len(data) * 0.8)

    # Split the data into training and testing sets
    train_data = data[:split_idx] # First 80% for training
    test_data = data[split_idx:]   # Last 20% for testing

    return train_data, test_data

# def splitDataSet(inputDataFrame, trainSetSize):

#     trainSet = inputDataFrame.sample(frac = trainSetSize)
#     testSet = inputDataFrame.drop(trainSet.index)
#     return trainSet, testSet
```

One Hot encoding and Data Splitting

```
In [423]: #One hot encoding
df['Y'] = df['Y'].map({'M':'0', 'B':'1'})
df['Y'].astype(float)

train_data, test_data = train_test_split(df)

Xtrain = train_data.drop('Y', axis=1)
Ytrain = train_data['Y'].astype(float)
Xtrainn = (Xtrain - Xtrain.mean())/Xtrain.std()

Xtests = test_data.drop('Y', axis=1)
Ytests = test_data['Y'].astype(float)
Xtestss = (Xtests - Xtests.mean())/Xtests.std()
```

```
In [ ]:
```

1A. Extending the Logistic Regression Class to include Stochastic Gradient Descent Method

```
In [424]: import numpy as np

class Loss:
    def __init__(self, x, y):
        self.x = x
        self.y = y

    def mean_square_loss(self, theta):
        predictions = np.dot(self.x, theta)
        return np.mean((self.y - predictions) ** 2)

    def mean_square_gradient(self, theta):
        predictions = np.dot(self.x, theta)
        gradient = -2 * np.dot(self.x.T, (self.y - predictions)) / len(self.y)
```

```

        return gradient

    def cross_entropy_loss(self, theta):
        predictions = self.sigmoid(np.dot(self.x, theta))
        return -np.mean(self.y * np.log(predictions) + (1 - self.y) * np.log(1 - predictions))

    def cross_entropy_gradient(self, theta):
        predictions = self.sigmoid(np.dot(self.x, theta))
        gradient = np.dot(self.x.T, (predictions - self.y)) / len(self.y)
        return gradient

    @staticmethod
    def sigmoid(z):
        return 1 / (1 + np.exp(-z))

class Optimization:
    def __init__(self, x, y):
        self.x = x
        self.y = y
        self.loss_obj = Loss(x, y)

    def mean_square_hessian(self):
        return 2 * np.dot(self.x.T, self.x) / len(self.x)

    def cross_entropy_hessian(self, theta):
        p = self.sigmoid(np.dot(self.x, theta))
        W = np.diag(p * (1 - p))
        hess = np.dot(self.x.T, np.dot(W, self.x))
        return hess

    @staticmethod
    def sigmoid(z):
        return 1 / (1 + np.exp(-z))

    def newtons_method(self, theta, lr=0.01, epochs=50, loss_trajectory=None):
        for epoch in range(epochs):
            gradient = self.loss_obj.cross_entropy_gradient(theta)
            hessian = self.cross_entropy_hessian(theta)
            inv_hessian = np.linalg.inv(hessian)
            theta -= lr * np.dot(inv_hessian, gradient)

            # Calculate loss for monitoring
            loss = self.loss_obj.cross_entropy_loss(theta)
            loss_trajectory.append(loss)
        return theta

    def sgd(self, theta, lr=0.01, epochs=50, batch_size=20, loss_trajectory=None):
        n_samples = self.x.shape[0]
        for epoch in range(epochs):

            indices = np.arange(n_samples)
            np.random.shuffle(indices)

            for i in range(0, n_samples, batch_size):
                batch_indices = indices[i:i + batch_size]
                x_batch = self.x[batch_indices]
                y_batch = self.y[batch_indices]

                predictions = Loss.sigmoid(np.dot(x_batch, theta))
                gradient = np.dot(x_batch.T, (predictions - y_batch)) / batch_size
                theta -= lr * gradient
                #print('Updated theta:', theta)

            # Calculate loss for monitoring
            loss = self.loss_obj.cross_entropy_loss(theta)
            loss_trajectory.append(loss)
        return theta

class LogisticRegression:
    def __init__(self):
        self.theta = None
        self.train_loss_trajectory = []

    def fit(self, x, y, lr=0.01, epochs=200, method="sgd", batch_size=1):
        self.theta = np.zeros(x.shape[1]) # Initialize coefficients (theta)
        optimizer = Optimization(x, y)

        if method == "sgd":
            self.theta = optimizer.sgd(self.theta, lr=lr, epochs=epochs, batch_size=batch_size, loss_trajectory=
        elif method == "newton":
            self.theta = optimizer.newtons_method(self.theta, lr=lr, epochs=epochs, loss_trajectory=self.train_

```

```

    return self.theta

def predict_proba(self, x):
    z = np.dot(x, self.theta)
    return Loss.sigmoid(z)

def predict(self, x, threshold=0.5):
    return (self.predict_proba(x) >= threshold).astype(int)


if __name__ == "__main__":

    model = LogisticRegression()

    Xtrainns = np.array(Xtrainn) # Convert Xtrainn to a NumPy array 'cus it's a DataFrame
    Ytrains = np.array(Ytrain)

    model.fit(Xtrainns, Ytrains, lr=0.01, epochs=300, method="sgd", batch_size=1)
    predictions = model.predict(Xtrainns)
    plt.plot(model.train_loss_trajectory, label='Train Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Cross Entropy Loss')
    plt.legend()
    plt.title("Stochastic Gradient Descent")
    plt.show()

```



In []:

B. Extending the Loss Class to include L1/L2 Regularization

```

In [437...] class Loss:
    def __init__(self, x, y, lambda_=0.01):
        self.x = x
        self.y = y
        self.lambda_ = lambda_ ## Adding lambda

    def mean_square_loss(self, theta):
        predictions = np.dot(self.x, theta)
        return np.mean((self.y - predictions) ** 2)

    def mean_square_gradient(self, theta):
        predictions = np.dot(self.x, theta)
        gradient = -2 * np.dot(self.x.T, (self.y - predictions)) / len(self.y)
        return gradient

    def cross_entropy_loss(self, theta):
        predictions = self.sigmoid(np.dot(self.x, theta))
        return -np.mean(self.y * np.log(predictions) + (1 - self.y) * np.log(1 - predictions))

    def L2_regularized_cross_entropy_loss(self, theta):
        predictions = self.sigmoid(np.dot(self.x, theta))
        loss = -np.mean(self.y * np.log(predictions) + (1 - self.y) * np.log(1 - predictions))

```

```

        loss -= (self.lambda_) * np.sum(theta ** 2)  ### Adding the L2 regularization
    return loss

def L1_regularized_cross_entropy_loss(self, theta):
    predictions = self.sigmoid(np.dot(self.x, theta))
    loss = -np.mean(self.y * np.log(predictions) + (1 - self.y) * np.log(1 - predictions))
    loss -= (self.lambda_) * np.sum(np.abs(theta))  ### Adding L1 Regularization
    return loss

def cross_entropy_gradient(self, theta):
    predictions = self.sigmoid(np.dot(self.x, theta))
    gradient = np.dot(self.x.T, (predictions - self.y)) / len(self.y)
    return gradient

def L2_regularized_cross_entropy_gradient(self, theta):
    predictions = self.sigmoid(np.dot(self.x, theta))
    gradient = np.dot(self.x.T, (predictions - self.y)) / len(self.y)
    gradient -= 2*(self.lambda_ * theta)

def L1_regularized_cross_entropy_gradient(self, theta):
    predictions = self.sigmoid(np.dot(self.x, theta))
    gradient = np.dot(self.x.T, (predictions - self.y)) / len(self.y)
    gradient -= self.lambda_ * np.sign(theta)

@staticmethod
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

```

In []:

C and E. Fitting the Logistic Regression Models with the combinations as stated and Generating the trajectory, accuracy report

In []:

```

In [440...] class Optimization:
    def __init__(self, x, y):
        self.x = x
        self.y = y
        self.loss_obj = Loss(x, y)

    def mean_square_hessian(self):
        return 2 * np.dot(self.x.T, self.x) / len(self.x)

    def cross_entropy_hessian(self, theta):
        p = self.sigmoid(np.dot(self.x, theta))
        W = np.diag(p * (1 - p))
        hess = np.dot(self.x.T, np.dot(W, self.x))
        return hess

    @staticmethod
    def sigmoid(z):
        return 1 / (1 + np.exp(-z))

    def newtons_method(self, theta, lr=0.01, epochs=50, loss_trajectory=None):
        for epoch in range(epochs):
            gradient = self.loss_obj.cross_entropy_gradient(theta)
            hessian = self.cross_entropy_hessian(theta)
            inv_hessian = np.linalg.inv(hessian)
            theta -= lr * np.dot(inv_hessian, gradient)

            # Calculate loss for monitoring
            loss = self.loss_obj.cross_entropy_loss(theta)
            loss_trajectory.append(loss)
        return theta

    def sgd(self, theta, lr=0.01, epochs=50, batch_size=20, loss_trajectory=None, loss_trajectory1=None, loss_tra
        n_samples = self.x.shape[0]
        for epoch in range(epochs):

            indices = np.arange(n_samples)
            np.random.shuffle(indices)

            # for i in range(0, n_samples, batch_size):
            batch_indices = indices[i:i + batch_size]
            x_batch = self.x[batch_indices]
            y_batch = self.y[batch_indices]

            predictions = Loss.sigmoid(np.dot(x_batch, theta))
            gradient = np.dot(x_batch.T, (predictions - y_batch)) / batch_size
            theta -= lr * gradient

```

```

        #print('Updated theta:', theta)

    # Calculate loss for monitoring
    loss = self.loss_obj.cross_entropy_loss(theta)
    loss1 = self.loss_obj.L2_regularized_cross_entropy_loss(theta)
    loss2 = self.loss_obj.L1_regularized_cross_entropy_loss(theta)

    loss_trajectory.append(loss)
    loss_trajectory1.append(loss1)
    loss_trajectory2.append(loss2)

    return theta

class LogisticRegression:
    def __init__(self):
        self.theta = None
        self.train_loss_trajectory = []
        self.train_loss_trajectory1 = []
        self.train_loss_trajectory2 = []

    def fit(self, x, y, lr=0.01, epochs=200, method="sgd", batch_size=1):
        self.theta = np.zeros(x.shape[1]) # Initialize coefficients (theta)
        optimizer = Optimization(x, y)

        if method == "sgd":
            self.theta = optimizer.sgd(self.theta, lr=lr, epochs=epochs, batch_size=batch_size, loss_trajectory=self.train_loss_trajectory,
                                       loss_trajectory1=self.train_loss_trajectory1, loss_trajectory2=self.train_loss_trajectory2)
        elif method == "newton":
            self.theta = optimizer.newtons_method(self.theta, lr=lr, epochs=epochs, loss_trajectory=self.train_loss_trajectory)
        return self.theta

    def predict_proba(self, x):
        z = np.dot(x, self.theta)
        return Loss.sigmoid(z)

    def predict(self, x, threshold=0.5):
        return (self.predict_proba(x) >= threshold).astype(int)



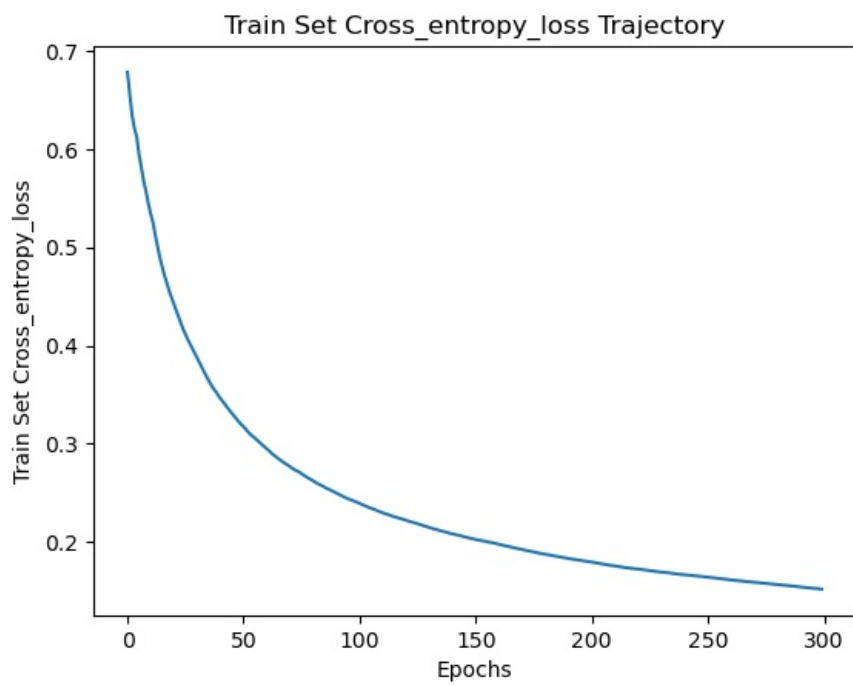
if __name__ == "__main__":
    model = LogisticRegression()

    Xtrainns = np.array(Xtrainn) # Convert Xtrainn to a NumPy array if it's a DataFrame
    Ytrains = np.array(Ytrain)

    model.fit(Xtrainns, Ytrains, lr=0.01, epochs=300, method="sgd", batch_size=20)
    predictions = model.predict(Xtrainns)
    loss_trajectories = [model.train_loss_trajectory, model.train_loss_trajectory1, model.train_loss_trajectory2]
    labels = ['Train Set Cross_entropy_loss', 'Train Set L2_regularized_cross_entropy_loss', 'Train Set L1_regularized_cross_entropy_loss']

    for i, trajectory in enumerate(loss_trajectories):
        plt.figure()
        plt.plot(loss_trajectories[i], label=labels[i])
        plt.xlabel('Epochs')
        plt.ylabel(labels[i])
        plt.title(f"{labels[i]} Trajectory")
    # plt.legend()
    plt.show()

```



In [441]: `if __name__ == "__main__":`

```

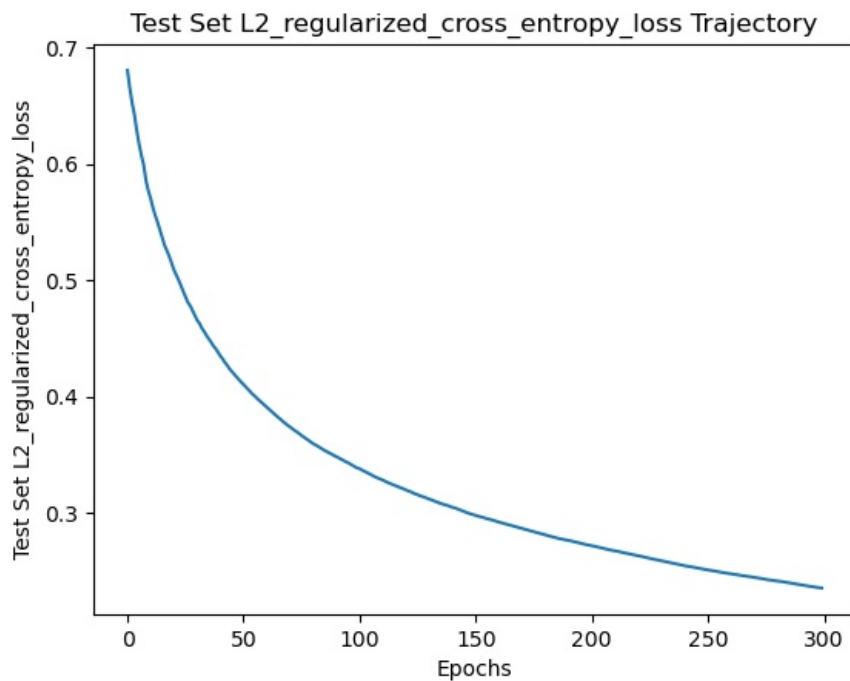

model2 = LogisticRegression()


Xtestss = np.array(Xtestss) # Convert Xtrainn to a NumPy array if it's a DataFrame
Ytests = np.array(Ytests)

model2.fit(Xtestss, Ytests, lr=0.01, epochs=300, method="sgd", batch_size=20)
predictions2 = model2.predict(Xtestss)
loss_trajectories2 = [model2.train_loss_trajectory, model2.train_loss_trajectory1, model2.train_loss_trajectory2]
labels = ['Test Set Cross_entropy_loss', 'Test Set L2_regularized_cross_entropy_loss', 'Test Set L1_regularized_cross_entropy_loss']

for i, trajectory in enumerate(loss_trajectories2):
    plt.figure()
    plt.plot(loss_trajectories2[i], label=labels[i])
    plt.xlabel('Epochs')
    plt.ylabel(labels[i])
    plt.title(f"{labels[i]} Trajectory")
    # plt.legend()
    plt.show()

```





```
In [442]: accuracy = np.mean(predictions == Ytrains)
print("Accuracy for Training Dataset :", (accuracy*100))

accuracy2 = np.mean(predictions2 == Ytests)
print("Accuracy for Test Dataset:", (accuracy2*100))
```

Accuracy for Training Dataset : 97.14285714285714
Accuracy for Test Dataset: 95.6140350877193


In []:

Confusion Matrix

```
In [446]: cm_lr = confusion_matrix(Ytests, predictions2)
cm_lr
```

```
Out[446]: array([[26,  0],
               [ 5, 83]], dtype=int64)
```

```
In [445]: sns.heatmap(cm_lr, annot=True, fmt='d')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix - Logistic Regression')
plt.show()
```

In []:

True Positives (TP) = 26: The model correctly predicted 26 instances as positive. True Negatives (TN) = 83: The model correctly predicted 83 instances as negative. False Positives (FP) = 5: The model incorrectly predicted 5 instances as positive when they were actually negative. False Negatives (FN) = 0: The model did not miss any actual positive instances

Accuracy= $\frac{TP+TN}{FP+FN+TP+TN}$ = 95.6%

In []:

In []:

```
In [435]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

class Loss:
    def __init__(self, x, y, lambda_=0.01):
        self.x = x
        self.y = y
        self.lambda_ = lambda_  ## Adding lambda

    def cross_entropy_loss(self, theta):
        predictions = self.sigmoid(np.dot(self.x, theta))
        return -np.mean(self.y * np.log(predictions) + (1 - self.y) * np.log(1 - predictions))

    @staticmethod
    def sigmoid(z):
        return 1 / (1 + np.exp(-z))

class Optimization:
    def __init__(self, x, y):
        self.x = x
        self.y = y
        self.loss_obj = Loss(x, y)

    def sgd(self, theta, lr=0.01, epochs=50, batch_size=20, loss_trajectory=None):
        n_samples = self.x.shape[0]
        for epoch in range(epochs):
            indices = np.arange(n_samples)
            np.random.shuffle(indices)

            for i in range(0, n_samples, batch_size):
                batch_indices = indices[i:i + batch_size]
                x_batch = self.x[batch_indices]
                y_batch = self.y[batch_indices]

                predictions = Loss.sigmoid(np.dot(x_batch, theta))
                gradient = np.dot(x_batch.T, (predictions - y_batch)) / batch_size
                theta -= lr * gradient

            # Calculate loss for monitoring
            loss = self.loss_obj.cross_entropy_loss(theta)
            loss_trajectory.append(loss)
```

```

        return theta

class LogisticRegression:
    def __init__(self):
        self.theta = None
        self.train_loss_trajectory = []

    def fit(self, x, y, lr=0.01, epochs=200, batch_size=1):
        self.theta = np.zeros(x.shape[1]) # Initialize coefficients (theta)
        optimizer = Optimization(x, y)
        self.theta = optimizer.sgd(self.theta, lr=lr, epochs=epochs, batch_size=batch_size, loss_trajectory=self.train_loss_trajectory)
        return self.theta

    def predict_proba(self, x):
        z = np.dot(x, self.theta)
        return Loss.sigmoid(z)

    def predict(self, x, threshold=0.5):
        return (self.predict_proba(x) >= threshold).astype(int)

# AIC calculation function based on cross-entropy loss and number of parameters
def compute_aic(model, X, y):
    predictions = model.predict_proba(X)
    loss = Loss(X, y).cross_entropy_loss(model.theta)
    num_params = len(model.theta)
    return 2 * num_params + 2 * loss * len(y)
    #return -2 * loss + 2 * num_params

# Backward Feature Selection with AIC
def backward_feature_selection(X, y, min_features=1):
    selected_features = list(X.columns)
    best_aic = float('inf') # Initialize with a very high AIC value
    best_features = selected_features.copy()

    while len(selected_features) >= min_features:
        # Train model with the current set of selected features
        X_selected = X[selected_features]
        model = LogisticRegression()
        model.fit(X_selected.values, y, lr=0.01, epochs=300, batch_size=20)

        # Compute AIC for the model with the current features
        current_aic = compute_aic(model, X_selected.values, y)

        # Initialize variables to track best feature to remove
        worst_feature = None
        best_aic_with_removal = current_aic

        # Try removing each feature and calculate the resulting AIC
        for feature in selected_features:
            temp_features = [f for f in selected_features if f != feature]
            X_temp = X[temp_features]

            # Train model without this feature
            temp_model = LogisticRegression()
            temp_model.fit(X_temp.values, y, lr=0.01, epochs=300, batch_size=20)
            temp_aic = compute_aic(temp_model, X_temp.values, y)

            # Check if this removal results in a lower AIC
            if temp_aic < best_aic_with_removal:
                best_aic_with_removal = temp_aic
                worst_feature = feature

        # If removing the worst feature improves AIC, update selected features
        if worst_feature is not None and best_aic_with_removal < current_aic:
            selected_features.remove(worst_feature)
            best_aic = best_aic_with_removal
            best_features = selected_features.copy()
        else:
            # Stop if no improvement is possible
            break

    return best_features, best_aic

# Example usage
if __name__ == "__main__":
    Xtrainnz = Xtrainn.copy() # Convert Xtrainn to a NumPy array if it's a DataFrame
    Ytrainnz = np.array(Ytrain.copy())

    # Perform backward feature selection
    best_features, best_aic = backward_feature_selection(Xtrainnz, Ytrainnz, min_features=1)
    print("Selected Features:", best_features)

```

```
print("Best AIC:", best_aic)
```

```
Selected Features: ['X2', 'X11', 'X20', 'X21', 'X23', 'X25', 'X28']  
Best AIC: 80.33208544091798
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

2. K-Fold Cross Validation

```
In [271...] Xtestss = np.array(Xtestss) # Convert Xtest to a NumPy array if it's a DataFrame (Already did for training data)  
Ytests = np.array(Ytests)
```

```
In [135...] Xtrain = train_data.drop('Y', axis=1)  
Ytrain = train_data['Y'].astype(float)  
Xtrainn = (Xtrain - Xtrain.mean())/Xtrain.std()
```

```
Xtests = test_data.drop('Y', axis=1)  
Ytests = test_data['Y'].astype(float)  
Xtestss = (Xtests - Xtests.mean())/Xtests.std()
```

```
In [136...] y_train_bank=pd.DataFrame(Ytrain.values.reshape(-1,1))  
y_test_bank=pd.DataFrame(Ytests.values.reshape(-1,1))  
x_train_bank=pd.DataFrame(Xtrainn.values)  
x_test_bank=pd.DataFrame(Xtestss.values)
```

```
In [137...] print('x_train_bank:',x_train_bank.shape)  
print('x_test_bank:',x_test_bank.shape)  
print('y_train_bank:',y_train_bank.shape)  
print('y_test_bank:',y_test_bank.shape)
```

```
x_train_bank : (455, 30)  
x_test_bank : (114, 30)  
y_train_bank : (455, 1)  
y_test_bank : (114, 1)
```

```
In [ ]:
```

```
In [365...] def logistic_function(X, beta):  
    z = np.dot(X,beta)  
    return 1 / (1 + np.exp(-z))  
  
def log_likelihood(x, y, beta):  
    z = np.dot(x, beta)  
    log = np.sum( y*z - np.log(1 + np.exp(z)) )  
    return log  
  
#L2-regularized cross-entropy loss  
  
betas = lambda x,y,beta,alpha,lamda : beta-alpha*(-2*np.dot(x.T,y-logistic_function(x, beta))+(2*lamda)*beta)  
  
def stochastic_gradient_descent(x_train,y_train,alpha,epochs,lamda,x_test,y_test):  
    m_train,n_features = np.shape(x_train)  
    ini_alpha = alpha  
    beta_hat = np.random.random(n_features).reshape(-1,1)  
    logtrain = []  
    logtest = []  
    y_hat = logistic_function(x_train,beta_hat)  
  
    chunk_size = 20  
  
    for i in range(epochs):  
        for chunk in range(len(x_train)//chunk_size):  
            x_chunk = x_train[chunk*chunk_size:min((chunk+1)*chunk_size,len(x_train))]  
            y_chunk = y_train[chunk*chunk_size:min((chunk+1)*chunk_size,len(y_train))]  
  
            beta_hat = betas(x_chunk,y_chunk,beta_hat,alpha,lamda)  
  
            y_hat=logistic_function(x_train,beta_hat)  
  
            logtest.append(log_likelihood(x_test,y_test,beta_hat))  
            logtrain.append(log_likelihood(x_train, y_train,beta_hat))  
  
    return logtest,logtrain,beta_hat
```

```
In [379...] import math as Math
```

```

def gridsearch(alpha,lamda):
    comb=[]
    for i in range(0,len(alpha)):
        for k in range(0,len(lamda)):
            comb.append(dict([('alpha',alpha[i]),('lamda',lamda[k])]))
    return comb

def data_k_divide(data,k):
    k_size=math.floor(len(data)/k)
    k_data=[]
    c=0
    for i in range (0,k):
        data_set=pd.DataFrame(data.head(0))
        for j in range(i*k_size,(i*k_size)+k_size):
            #data_set=data_set.append(data.iloc[j])
            data_set=pd.concat([data_set, data.iloc[[j]]])
            c=c+1
        k_data.append(data_set)

    #adding datas which are remaining at the end of k division
    # for j in range(c,len(data)):
    #     k_data[k-1]=k_data[k-1].append(data.iloc[j])
    return k_data

def k_data_train_test(x,y,k):
    k_folded_data=[]
    for i in range(0,k):
        x_test=x[i]
        y_test=y[i]
        x_train=pd.DataFrame()
        y_train=pd.DataFrame()
        for j in range(0,k):
            if i!=j:

                x_train = pd.concat([x_train, x[j]])
                y_train = pd.concat([y_train, y[j]])

        final_data=dict([('x',x_train),('y',y_train),('xt',x_test),('yt',y_test)])
        k_folded_data.append(final_data)
    return k_folded_data

def kfold(x_train,y_train,k,x_test,y_test):
    x_train_k=data_k_divide(x_train,k)
    y_train_k=data_k_divide(y_train,k)
    data=k_data_train_test(x_train_k,y_train_k,k)
    return data

```

```

In [513.. import warnings

# Suppress FutureWarnings
warnings.simplefilter(action='ignore', category=FutureWarning)

alpha=[0.0001,0.00001,0.001,0.05,0.008]
lamda=[0.0000001,0.1,0.5,0.01,0.006]
epochs=20
k=3
parameter=gridsearch(alpha,lamda)
log_test=[]
log_last=[]
log_train=[]
alpha_com=[]
lamda_com=[]
avg_log=[]

## Just commenting to save space

for i in range (0,len(parameter)):
    k_folded_data=kfold(x_train_bank,y_train_bank,k,x_test_bank,y_test_bank)
    for j in range(0,k):
        logtest,logtrain,beta_hat =stochastic_gradient_descent(k_folded_data[j]['x'],k_folded_data[j]['y'],
        log_last.append(logtest[-1])
        log_test.append(logtest)
        log_train.append(logtrain)
    alpha_com.append(parameter[i]['alpha'])
    lamda_com.append(parameter[i]['lamda'])
    avg_log.append(np.mean(log_last))


```

```

In [384.. from mpl_toolkits import mplot3d
fig = plt.figure()
fig.set_figheight(10)
fig.set_figwidth(10)
ax = plt.axes(projection= '3d')

```

```
ax.scatter3D(alpha_com, lamda_com, avg_log, c=avg_log, cmap='hsv')
ax.set_xlabel('alpha')
ax.set_ylabel('lamda')
ax.set_zlabel('LOG')
plt.show()
```



```
In [397...] combined_df = pd.DataFrame({
    'alpha': alpha_com,
    'lamda': lamda_com,
    'avg_log': avg_log
})
combined_df.sort_values('avg_log', ascending=False)
```

Out [397...

| | alpha | lamda | avg_log |
|----|---------|--------------|-------------|
| 4 | 0.00010 | 6.000000e-03 | -84.380376 |
| 2 | 0.00010 | 5.000000e-01 | -84.741091 |
| 3 | 0.00010 | 1.000000e-02 | -85.143950 |
| 1 | 0.00010 | 1.000000e-01 | -88.523028 |
| 0 | 0.00010 | 1.000000e-07 | -96.002594 |
| 24 | 0.00800 | 6.000000e-03 | -220.392050 |
| 23 | 0.00800 | 1.000000e-02 | -229.112417 |
| 5 | 0.00001 | 1.000000e-07 | -233.649491 |
| 22 | 0.00800 | 5.000000e-01 | -238.570482 |
| 21 | 0.00800 | 1.000000e-01 | -248.757093 |
| 20 | 0.00800 | 1.000000e-07 | -260.047036 |
| 19 | 0.05000 | 6.000000e-03 | -272.495024 |
| 18 | 0.05000 | 1.000000e-02 | -286.111628 |
| 17 | 0.05000 | 5.000000e-01 | -301.280896 |
| 16 | 0.05000 | 1.000000e-01 | -318.123547 |
| 15 | 0.05000 | 1.000000e-07 | -337.281161 |
| 6 | 0.00001 | 1.000000e-01 | -339.223252 |
| 14 | 0.00100 | 6.000000e-03 | -358.811995 |
| 13 | 0.00100 | 1.000000e-02 | -383.350229 |
| 12 | 0.00100 | 5.000000e-01 | -411.505411 |
| 7 | 0.00001 | 5.000000e-01 | -419.611547 |
| 11 | 0.00100 | 1.000000e-01 | -444.288197 |
| 8 | 0.00001 | 1.000000e-02 | -473.424552 |
| 10 | 0.00100 | 1.000000e-07 | -483.293431 |
| 9 | 0.00001 | 6.000000e-03 | -530.000710 |

Using the best parameters alpha,lamda

In [398...

```
## We already generated the Optimal Betas in the codes above

#beta_hat3
def predict(x, threshold=0.5):
    return (x >= threshold).astype(int)

def predict_proba(x,beta):
    z = np.dot(x, beta)
    return 1 / (1 + np.exp(-z))

pre_values4 = predict_proba(x_test_bank,beta_hat)
pre_probability4 = predict(pre_values4)

### Now fitting it to our test data

accuracy4 = np.mean(pre_probability4 == y_test_bank)
print(" Final Accuracy for Test Dataset - Optimum hyper-parameters:", (accuracy4 * 100))
```

Final Accuracy for Test Dataset - Optimum hyper-parameters: 92.10526315789474

In [385...

```
# logtest1, logtrain1, beta_hat1 = stochastic_gradient_descent(
#     x_train_bank,
#     y_train_bank,
#     alpha=0.0001,
#     epochs=20,
#     lamda=0.0000001,
#     x_test=x_test_bank,
#     y_test=y_test_bank
# )
```

In [389...

In []:

3 Coordinate Descent for L1-Regularized Linear Regression

A. L1-Regularized Linear Regression

```
In [360.. df1 = pd.read_csv('regression2.csv')
df1.head()
```

```
Out[360..
```

| | x1 | x2 | x3 | x4 | x5 | x6 | x7 | x8 | x9 | y |
|---|-----------|-----------|----------|-----------|-----------|-----------|----------|-----------|----------|------------|
| 0 | 0.496714 | -0.138264 | 0.647689 | 1.523030 | -0.234153 | -0.068678 | 0.419500 | 0.998859 | 0.210385 | 10.042597 |
| 1 | -0.234137 | 1.579213 | 0.767435 | -0.469474 | 0.542560 | -0.369752 | 0.588956 | -0.452418 | 0.433443 | -6.907481 |
| 2 | -0.463418 | -0.465730 | 0.241962 | -1.913280 | -1.724918 | 0.215827 | 0.058546 | -0.941923 | 1.002438 | -11.705652 |
| 3 | -0.562288 | -1.012831 | 0.314247 | -0.908024 | -1.412304 | 0.569502 | 0.098751 | -0.788289 | 0.880582 | -4.990023 |
| 4 | 1.465649 | -0.225776 | 0.067528 | -1.424748 | -0.544383 | -0.330909 | 0.004560 | -0.989354 | 0.434624 | -2.982742 |

```
In [ ]:
```

```
In [399.. train_data2, test_data2 = train_test_split(df1)

Xtrain2 = train_data2.drop('y', axis=1)
Ytrain2 = train_data2['y']
Xtrainn2 = (Xtrain2 - Xtrain2.mean())/Xtrain2.std()
Ytrainn2 = (Ytrain2 - Ytrain2.mean())/Ytrain2.std()

Xtests2 = test_data2.drop('y', axis=1)
Ytests2 = test_data2['y'].astype(float)
Xtestss2 = (Xtests2 - Xtests2.mean())/Xtests2.std()
Ytestss2 = (Ytests2 - Ytests2.mean())/Ytests2.std()
```

```
In [453.. def mean_square_loss(x,y,theta,lambda_):
    predictions = np.dot(x, theta)
    loss = np.mean((y - predictions) ** 2)
    loss+= (lambda_) * np.sum(np.abs(theta))
    return loss

def mean_square_gradient(x,y,theta,lambda_):
    predictions = np.dot(x, theta)
    gradient = -2 * np.dot(x.T, (y - predictions)) / len(y)
    gradient+= lambda_ * np.sign(theta)
    return gradient

def sgd(x,y,theta, lr, epochs, batch_size, loss_trajectory):
    n_samples = x.shape[0]
    for epoch in range(epochs):

        indices = np.arange(n_samples)
        np.random.shuffle(indices)

        for i in range(0, n_samples, batch_size):
            batch_indices = indices[i:i + batch_size]
            x_batch = x[batch_indices]
            y_batch = y[batch_indices]

            predictions = np.dot(x_batch, theta)
            gradient = mean_square_gradient(x_batch,y_batch,theta,lambda_)
            theta -= lr * gradient
            #print('Updated theta:', theta)

            # Calculate loss for monitoring
            loss = mean_square_loss(x,y,theta,lambda_)

            loss_trajectory.append(loss)

        return theta,loss_trajectory


lambda_ = 0.5
Xtrainns2 = np.array(Xtrainn2) # Convert Xtrainn to a NumPy array if it's a DataFrame
Ytrainns2 = np.array(Ytrainn2)
theta = np.zeros(Xtrainns2.shape[1])
```

```
In [454.. ttheta,traj = sgd(Xtrainns2,Ytrainns2,theta, lr=0.01, epochs=150, batch_size=20, loss_trajectory=[])
predictionss = np.dot(Xtrain2, ttheta)
print('Optimum Theta after iterations are :', ttheta)
```

```
Optimum Theta after iterations are : [ 0.20284814 -0.05139114  0.01080433  0.48773802 -0.00337104  0.00531932
 0.00295894  0.03459513 -0.00196837]
```

```
In [455.. plt.plot(traj)
```

Out[455... [<matplotlib.lines.Line2D at 0x1f5a63df190>]



```
In [456... from sklearn.metrics import mean_squared_error, r2_score
r2 = r2_score(Ytrains2, predictionss)
print("R-squared (R2):", r2)
```

R-squared (R2): 0.7005693847824108

In []:

B. Coordinate Descent

In []:

```
In [511... def soft_threshold(beta, reg):

    if beta < reg:
        return beta + reg
    elif beta > reg:
        return beta - reg
    else:
        return 0

# Coordinate Descent algorithm for Lasso regression
def coordinate_descent(x, y, lambda_, epochs):
    m_train, n_features = np.shape(x)
    beta = np.zeros(n_features).reshape(-1, 1)
    beta_hist = np.zeros((epochs + 1, n_features)) # To track the evolution of coefficients

    # Coordinate descent loop
    for j in range(epochs):
        for i in range(n_features):

            # Choosing the feature (coordinate) to update
            _x = x[:, i]
            x_coor = np.delete(x, i, axis=1) # All features except i-th
            beta_coor = np.delete(beta, i, axis=0) # All coefficients except i-th

            # Compute the linear part of the update
            col = x_coor.dot(beta_coor)
            num = ((y - col).T).dot(_x) # Numerator
            den = _x.T.dot(_x) # Denominator
            update = num / den

            # Apply soft-thresholding for L1 regularization (Lasso)
            reg = lambda_ / (_x.T.dot(_x)) # Regularization term
            update = soft_threshold(np.mean(update), reg) # Apply soft-thresholding

            # the update should be scalar(I think..), and then update the coefficient
            beta[i] = update

        # Store the coefficients after each epoch
        beta_hist[j + 1] = beta.ravel()

    return beta, beta_hist

lambda_ = 0.7
```




```
epochs = 500

beta, beta_hist = coordinate_descent(Xtrainns2, Ytrains2, lambda_, epochs)

# The final coefficients are in `beta` while `beta_hist` tracks the evolution of the coefficients
print("Final coefficients:", beta)
#rint(beta_hist)
```

```
Final coefficients: [[0.00083667]
 [0.00088599]
 [0.00090127]
 [0.00065628]
 [0.00078307]
 [0.00093803]
 [0.00100089]
 [0.00036851]
 [0.00083157]]
```

```
In [512... plt.plot(beta_hist)
plt.xlabel('Epochs')
plt.ylabel('Coefficient change')
plt.show()
```



In []:

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js