

```
In [285]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from mpl_toolkits.mplot3d import Axes3D
```

```
In [ ]:
```

1. KNN Imputation and Classification

```
In [ ]:
```

a.

```
In [92]: df1 = pd.read_csv('DodgerLoopGame\DodgerLoopGame_TRAIN.txt', sep="\s+", header=None)
df2 = pd.read_csv('DodgerLoopGame\DodgerLoopGame_TEST.txt', sep="\s+", header=None)
df2.head()
```

```
Out[92]:
```

	0	1	2	3	4	5	6	7	8	9	...	279	280	281	282	283	284	285	286	287	288
0	1.0	2.0	8.0	7.0	6.0	2.0	8.0	5.0	10.0	3.0	...	11.0	7.0	12.0	4.0	9.0	13.0	3.0	5.0	5.0	3.0
1	1.0	8.0	5.0	10.0	11.0	9.0	10.0	7.0	18.0	11.0	...	13.0	7.0	6.0	7.0	1.0	2.0	6.0	7.0	8.0	6.0
2	1.0	1.0	9.0	5.0	2.0	5.0	5.0	5.0	13.0	6.0	...	10.0	5.0	14.0	9.0	13.0	10.0	9.0	5.0	11.0	6.0
3	1.0	2.0	6.0	6.0	3.0	8.0	6.0	3.0	4.0	3.0	...	8.0	13.0	9.0	11.0	10.0	4.0	6.0	6.0	3.0	10.0
4	1.0	8.0	6.0	5.0	11.0	7.0	4.0	10.0	8.0	8.0	...	19.0	13.0	15.0	17.0	10.0	8.0	12.0	16.0	15.0	13.0

5 rows × 289 columns

```
In [96]: NaTrain = df1.isna().sum().sum()
NaTest = df2.isna().sum().sum()
print('Numberof NaN values in Train Datasets are',NaTrain)
print('Numberof NaN values in Test Datasets are',NaTest)
```

Numberof NaN values in Train Datasets are 65
Numberof NaN values in Test Datasets are 272

```
In [ ]:
```

```
In [ ]:
```

b. KNN Imputer

```
In [98]: from sklearn.impute import KNNImputer
from scipy.spatial.distance import cdist

impute = KNNImputer(n_neighbors=3)

# Train Dataset
df1_imput = impute.fit_transform(df1)
df1_imputed = pd.DataFrame(df1_imput, columns=df1.columns)

# Test Dataset
df2_imput = impute.fit_transform(df2)
df2_imputed = pd.DataFrame(df2_imput, columns=df2.columns)

print("\nTrain Dataset Number of NaN values after Imputation:", df1_imputed.isna().sum().sum())
print("\nTest Dataset Number of NaN values after Imputation:", df2_imputed.isna().sum().sum())

#mse = mean_squared_error(df2_imput, df2_imputed)
#print(mse)
```

Train Dataset Number of NaN values after Imputation: 0

Test Dataset Number of NaN values after Imputation: 0

```
In [146]: from sklearn.impute import KNNImputer
from scipy.spatial.distance import cdist

def knn_imputer_grid_search(X_train, X_test, k_values):
    best_k = 1
    best_mean_distance = float('inf')

    for k in k_values:
```

```

imputer = KNNImputer(n_neighbors=k)
train_imputed = imputer.fit_transform(X_train)
test_imputed = imputer.transform(X_test)
# Checking shapes to see if it aligns or not ..
# print("Shape of imputed values after flattening:", np.isnan(imputed_values).sum())

# distances = cdist(original_values_without_nan.reshape(-1, 1), imputed_values.reshape(-1, 1), metric='euclidean')
distances = cdist(train_imputed, test_imputed, metric='euclidean') # calculating the pairwise Euclidean
mean_distance = np.mean(distances) # Getting the mean distances

print(f"K={k}, Mean Distance={mean_distance:.4f}")

# Update the best K if current mean distance is lower
if mean_distance < best_mean_distance:
    best_mean_distance = mean_distance
    best_k = k

print(f"Optimal number of neighbors (K): {best_k}") # Using the best K to impute both train and test data
print(f"Best Mean Distance: {best_mean_distance:.4f}")

# Final imputation with the best K
final_imputer = KNNImputer(n_neighbors=best_k)
# Using the best k for the input data
train_imputed = pd.DataFrame(final_imputer.fit_transform(X_train), columns=X_train.columns)
# test_imputed = pd.DataFrame(final_imputer.transform(X_test), columns=X_test.columns) # To perform imputation

return best_k

k_values = [1, 3, 5, 7, 9]
optimal_k = knn_imputer_grid_search(df1, df2, k_values)
print(f"Optimal K: {optimal_k}")

```

K=1, Mean Distance=179.2819
K=3, Mean Distance=178.9790
K=5, Mean Distance=178.8261
K=7, Mean Distance=178.8077
K=9, Mean Distance=178.7810
Optimal number of neighbors (K): 9
Best Mean Distance: 178.7810
Optimal K: 9

```

In [176]: def knn_imputer_grid_search(X_train, k_values):
# #     best_k = 1
# #     best_mean_distance = float('inf')

# #     missing_mask = ~X_train.isna()

# #     for k in k_values:

# #         imputer = KNNImputer(n_neighbors=k)

# #         train_imputed = imputer.fit_transform(X_train)

# #         # Getting the original and imputed values where the original values are missing
# #         original_values = X_train[missing_mask].values.ravel()
# #         imputed_values = train_imputed[missing_mask].ravel()

# #         original_values_without_nan = original_values[~np.isnan(original_values)]
# #         print("Shape of original values after flattening:", original_values_without_nan.shape)
# #         print("Shape of imputed values after flattening:", imputed_values.shape)

# #         # Checking shapes to see if it aligns or not ..
# #         # print("Shape of original values after dropping 65:", np.isnan(original_values_dropped).sum())
# #         # print("Shape of imputed values after flattening:", np.isnan(imputed_values).sum())

# #         # Compute pairwise Euclidean distance between original and imputed values
# #         distances = cdist(original_values_without_nan.reshape(-1, 1), imputed_values.reshape(-1, 1), metric='euclidean')
# #         #distances = cdist(X_train.values.reshape(-1, 1), imputed_values.reshape(-1, 1), metric='euclidean')

# #         # Compute the mean distance
# #         mean_distance = np.mean(distances)

# #         print(f"K={k}, Mean Distance={mean_distance:.4f}")

# #         # Update the best K if current mean distance is lower
# #         if mean_distance < best_mean_distance:
# #             best_mean_distance = mean_distance
# #             best_k = k

# #     return best_k

```

```

## k values = [1, 3, 5, 7, 9]
## optimal_k = knn_imputer_grid_search(df1, k_values)
## print(f"Optimal K: {optimal_k}")

```

```

# OUTPUT

```

```

# Shape of original values after flattening: (5715,)
# Shape of imputed values after flattening: (5715,)
# K=1, Mean Distance=14.9319
# Shape of original values after flattening: (5715,)
# Shape of imputed values after flattening: (5715,)
# K=3, Mean Distance=14.9319
# Shape of original values after flattening: (5715,)
# Shape of imputed values after flattening: (5715,)
# K=5, Mean Distance=14.9319
# Shape of original values after flattening: (5715,)
# Shape of imputed values after flattening: (5715,)
# K=7, Mean Distance=14.9319
# Shape of original values after flattening: (5715,)
# Shape of imputed values after flattening: (5715,)
# K=9, Mean Distance=14.9319
# Optimal number of neighbors (K): 1
# Best Mean Distance: 14.9319
# Optimal K: 1

```

```

In [ ]:

```

c. KNN Classifier

```

In [ ]:

```

```

In [273] def train_val_test_split(X,y):
# Calculate the split indices
split_train_idx = int(len(data) * 0.7) # 70% for training
split_temp_idx = int(len(data) * 0.85) # 85% for training + validation (so 15% remains for testing)

# Split the data into training, validation, and test sets
train_data = data[:split_train_idx] # First 70% for training
validation_data = data[split_train_idx:split_temp_idx] # Next 15% for validation
test_data = data[split_temp_idx:] # Last 15% for testing

train_data = data[:split_train_idx] # First 70% for training
validation_data = data[split_train_idx:split_temp_idx] # Next 15% for validation
test_data = data[split_temp_idx:]

return train_data, validation_data, test_data

def train_test_split(data, test_size=0.2):

split_idx = int(len(data) * (1 - test_size))

# Split the data into train and test sets
train_data = data[:split_idx]
test_data = data[split_idx:]

return train_data, test_data

```

```

In [ ]:

```

```

In [ ]:

```

```

In [178] from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split

def EuclideanDistance(a, b):
return np.sqrt(np.sum((a - b) ** 2))

def y_prediction(k_neighbors, y_train):
y_pred = [y_train[i] for i in k_neighbors]
return Counter(y_pred).most_common(1)[0][0] # Return the most common label (majority voting)

def predict_knn_class(x_train, y_train, k, z): # KNN for classification
distances = []

for i in range(len(x_train)): # We need to Calculate the Euclidean distance from the test point (z) to all

```

```

        distance = EuclideanDistance(z, x_train[i])
        distances.append((i, distance)) # Storing the index and corresponding distance

    distances.sort(key=lambda x: x[1])
    k_neighbors = [distances[i][0] for i in range(k)] # taking indices of the k nearest neighbors

    return y_prediction(k_neighbors, y_train) # majority voting prediction

def grid_search_knn(X, y, k_values): # Function to perform grid search for finding optimal K value

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    best_k = None
    best_accuracy = 0

    for k in k_values: # Trying different values of K and checking
        y_train_pred = []
        for i in range(len(X_train)):
            y_train_pred.append(predict_knn_class(X_train, y_train, k, X_train[i])) # Get predictions for train

        accuracy_train = accuracy_score(y_train, y_train_pred)

        y_test_pred = [] # Predict for the test set
        for i in range(len(X_test)):
            y_test_pred.append(predict_knn_class(X_train, y_train, k, X_test[i])) # Getting test set prediction

        accuracy_test = accuracy_score(y_test, y_test_pred)

        if accuracy_test > best_accuracy:
            best_accuracy = accuracy_test
            best_k = k

    print(f"K={k}, Train Accuracy: {accuracy_train:.4f}, Test Accuracy: {accuracy_test:.4f}")

    print(f"Optimal K: {best_k}")
    print(f"Best Test Accuracy with K={best_k}: {best_accuracy:.4f}")

    return best_k, best_accuracy

def main_knn_example():

    # Creating dataset (500 rsamples, 5 features) ## There were no target columns in the dodgr loop dataset
    np.random.seed(42)
    X = np.random.rand(500, 5) #
    y = np.random.randint(0, 2, 500) # 0 or 1

    k_values = [1, 3, 5, 7, 9] # Perform grid search over different K values
    grid_search_knn(X, y, k_values)

main_knn_example()

```

```

K=1, Train Accuracy: 1.0000, Test Accuracy: 0.5100
K=3, Train Accuracy: 0.7750, Test Accuracy: 0.4500
K=5, Train Accuracy: 0.6875, Test Accuracy: 0.5100
K=7, Train Accuracy: 0.6175, Test Accuracy: 0.4800
K=9, Train Accuracy: 0.5800, Test Accuracy: 0.5400
Optimal K: 9
Best Test Accuracy with K=9: 0.5400

```

In [171]: df1.head(3)

```

Out[171]:
   0    1    2    3    4    5    6    7    8    9  ...  279  280  281  282  283  284  285  286  287  288
0  1.0  7.0  3.0  6.0  11.0  8.0  6.0  6.0  10.0  4.0  ...  12.0  5.0  9.0  4.0  4.0  6.0  9.0  5.0  16.0  8.0
1  1.0  9.0  10.0  5.0  7.0  10.0  9.0  5.0  6.0  8.0  ...   8.0  5.0  4.0  8.0  6.0  11.0  5.0  8.0  9.0  6.0
2  1.0  12.0  18.0  11.0  11.0  19.0  17.0  4.0  6.0  8.0  ...  10.0  9.0  11.0  8.0  4.0  7.0  3.0  6.0  3.0  6.0

```

3 rows × 289 columns

In [169]: ## df1.iloc[1] #better for numpy

```
Out[169... 0      1.0
1      9.0
2     10.0
3      5.0
4      7.0
...
284    11.0
285     5.0
286     8.0
287     9.0
288     6.0
Name: 1, Length: 289, dtype: float64
```

In []:

2. Decision Trees

```
In [286... df3 = pd.read_csv('iris\iris.data', header=None)
df3[4].value_counts()
```

```
Out[286... 4
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
Name: count, dtype: int64
```

```
In [192... # df4 = pd.get_dummies(df3, columns=[4], drop_first=True)
# df4.head()
```

```
Out[192...    0    1    2    3  4_Iris-versicolor  4_Iris-virginica
0  5.1  3.5  1.4  0.2                False                False
1  4.9  3.0  1.4  0.2                False                False
2  4.7  3.2  1.3  0.2                False                False
3  4.6  3.1  1.5  0.2                False                False
4  5.0  3.6  1.4  0.2                False                False
```

```
In [287... df3[4] = df3[4].map({'Iris-setosa':'0', 'Iris-versicolor':'1', 'Iris-virginica':'2'})
df3[4] = df3[4].astype(float)
df3.head()
```

```
Out[287...    0    1    2    3    4
0  5.1  3.5  1.4  0.2  0.0
1  4.9  3.0  1.4  0.2  0.0
2  4.7  3.2  1.3  0.2  0.0
3  4.6  3.1  1.5  0.2  0.0
4  5.0  3.6  1.4  0.2  0.0
```

```
In [288... x = df3.drop(columns=4)
y = df3[4]
```

In []:

a) Class for Tree Node

```
In [289... class Node:
    def __init__(self, feature=None, threshold=None, left=None, right=None, *, value=None):
        self.feature = feature
        self.threshold = threshold
        self.left = left
        self.right = right
        self.value = value # Predicted value for leaf nodes

    def is_leaf_node(self):
        return self.value is not None
```

b) Functions to build the Tree using RSS as the criterion

```
In [290... def rss(y):
    if len(y) == 0: # Handle empty data
        return 0
    mean_y = np.mean(y)
```

```
return np.sum((y - mean_y) ** 2)
```

```
class DecisionTree:
```

```
def __init__(self, min_samples_split=2, max_depth=100, n_feats=None):
    self.min_samples_split = min_samples_split
    self.max_depth = max_depth
    self.n_feats = n_feats
    self.root = None
```

```
def fit(self, X, y):
    self.n_feats = X.shape[1]
    self.root = self._grow_tree(X, y)
```

```
# def predict(self, X):
#     return np.array([self._traverse_tree(x, self.root) for x in X])
```

```
def _grow_tree(self, X, y, depth=0):
    n_samples, n_features = X.shape
```

```
    # Stopping criteria
```

```
    if (depth >= self.max_depth or n_samples < self.min_samples_split):
        leaf_value = np.mean(y)
        return Node(value=leaf_value)
```

```
    feat_idx = np.random.choice(n_features, self.n_feats, replace=False)
```

```
    # Greedy search
```

```
    best_feat, best_thresh = self._best_criteria(X, y, feat_idx)
```

```
    # Growing the children
```

```
    left_idx, right_idx = self._split(X[:, best_feat], best_thresh)
    left = self._grow_tree(X[left_idx, :], y[left_idx], depth + 1)
    right = self._grow_tree(X[right_idx, :], y[right_idx], depth + 1)
    return Node(best_feat, best_thresh, left, right)
```

```
def _best_criteria(self, X, y, feat_idx):
```

```
    best_gain = -1
    split_idx, split_thresh = None, None
    for feat_idx in feat_idx:
        X_column = X[:, feat_idx]
        thresholds = np.unique(X_column)
        for threshold in thresholds:
            gain = self._calculate_rss(y, X_column, threshold)

            if gain > best_gain:
                best_gain = gain
                split_idx = feat_idx
                split_thresh = threshold
```

```
    return split_idx, split_thresh
```

```
def _calculate_rss(self, y, X_column, split_thresh):
```

```
    left_idx, right_idx = self._split(X_column, split_thresh)
```

```
    # If no split (empty subset), return infinity to avoid it
```

```
    if len(left_idx) == 0 or len(right_idx) == 0:
        return float("inf")
```

```
    rss_left = rss(y[left_idx])
    rss_right = rss(y[right_idx])
```

```
    return rss_left + rss_right
```

```
def _split(self, X_column, split_thresh):
```

```
    left_idx = np.argwhere(X_column <= split_thresh).flatten()
    right_idx = np.argwhere(X_column > split_thresh).flatten()
```

```
    if len(left_idx) == 0 or len(right_idx) == 0:
        print(f"Warning: One of the splits is empty (threshold: {split_thresh}).")
    return left_idx, right_idx
```

```
# def _traverse_tree(self, x, node):
```

```
#     if node.is_leaf_node():
#         if node.value is None:
#             print("error: leaf node value is None")
#             raise ValueError("leaf node has no value.")
#         return node.value
```

```

#     if node.feature is None or node.threshold is None:
#         print("Error: Node feature or threshold is None")
#         raise ValueError("Node feature or threshold is None during traversal.")

#     print(f"Traversing node: feature {node.feature}, threshold {node.threshold}, value {node.value}")
#     if x[node.feature] <= node.threshold:
#         return self._traverse_tree(x, node.left)
#     return self._traverse_tree(x, node.right)

def mean_squared_error(self, y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)

# x = np.array(x)
# y = np.array(y).flatten()
# X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
# clf = DecisionTree(max_depth=10)
# clf.fit(X_train, y_train)

# # Make predictions on the test set
# y_pred = clf.predict(X_test)

# mse = clf.mean_squared_error(y_test, y_pred)
# print("Mean Squared Error:", mse)

```

c) Transversing the tree and making predictions

```

In [ ]: def _traverse_tree(self, x, node):
        if node.is_leaf_node():
            return node.value

        if x[node.feature] <= node.threshold:
            return self._traverse_tree(x, node.left)
        return self._traverse_tree(x, node.right)

def predict(self, X):
    return np.array([self._traverse_tree(x, self.root) for x in X])

```

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js