

Lab Course Machine Learning

Exercise Sheet 6

Prof. Dr. Dr. Lars Schmidt-Thieme
Jung Min Choi

HiWi: Harish Malik
Submission deadline : November 28, 2024

General Instructions

1. Data should be normalized.
2. Train to Test split should be 80-20 / with Validation 70-15-15
3. Convert any non-numeric values to numeric values. For example you can replace a country name with an integer value or more appropriately use one-hot encoding.

1 Naive Bayesian Classifier

(10 points)

In this assignment, you will implement a Naive Bayesian Classifier from scratch and compare its performance with the implementation from `sklearn`.

1. **[6 Points]** Implement a Naive Bayesian Classifier from scratch. Use the provided dataset `logistic.csv`. Load the data and split it into training and test sets using `train_test_split` from `sklearn`. Feature scaling can also be applied using `sklearn` methods if you think it necessary.
2. **[2 Points]** Compute the following evaluation metrics. (do not use `sklearn` or any other library for these computations):
 - Accuracy
 - Precision
 - Recall
 - F1-Score
3. **[2 Points]** Train and evaluate the Naive Bayesian Classifier implemented using `sklearn`'s `GaussianNB`. Compare the evaluation metrics (accuracy, precision, recall, F1-score) between your custom implementation and the `sklearn` implementation.

2 MLP Regressor with PyTorch

(10 points)

You will build and train a simple Multilayer Perceptron (MLP) regressor using the California Housing dataset.

1. **[1 Point]** Download the California Housing dataset using `sklearn`. Perform the following preprocessing steps:
 - Split the dataset into training, validation, and test sets.
 - Apply feature scaling (e.g., normalization or standardization)
 - You can use `sklearn`.
2. **[6 Points]** Build and train a Multilayer Perceptron (MLP) regressor using `PyTorch`. The MLP should satisfy the following requirements:
 - The architecture should allow customization of:
 - The number of hidden layers (n).
 - The activation function (e.g., `ReLU`, `Tanh`, etc.).

- Use the Adam optimizer and Mean Squared Error (MSE) as the loss function.
 - Use validation dataset for training.
 - You can use torch for the optimizer and the loss function.
3. [1 Point] Evaluate the model on the test dataset using MSE.
 4. [2 Point] Plot the training, and validation loss trajectories.

Hints:**1. PyTorch Tutorials and Installation:**

- Students who are unfamiliar with PyTorch can explore the official tutorials at the following link: <https://pytorch.org/tutorials/>.
- To install PyTorch, refer to the installation guide available here: <https://pytorch.org/get-started/locally/>.

2. Pseudocode for MLP Regressor:

- Define the MLP architecture:
 - Create a class that inherits from `torch.nn.Module`.
 - Define the layers of the MLP (input, hidden, and output) using `torch.nn.Linear`.
 - Use an activation function (e.g., `torch.nn.ReLU`) after each hidden layer.
- Implement the forward pass in the class.
- Preprocess the data:
 - Split the dataset into training, validation, and test sets.
 - Normalize or standardize the features.
 - Convert the data to PyTorch tensors.
 - Use `torch.utils.data.DataLoader` to create batches for training and validation.
- Define the optimizer (e.g., `torch.optim.Adam`) and loss function (e.g., `torch.nn.MSELoss`).
- Train the model:
 - Loop over a specified number of epochs.
 - For each epoch:
 - * Perform forward and backward passes for training data.
 - * Compute and track losses for training and validation datasets.
- Evaluate the model on the test set and compute MSE and RMSE.

3. PyTorch Functions:

- **Data Handling:**
 - `torch.utils.data.DataLoader` for creating data batches.
 - `torch.tensor` to convert numpy arrays or data into PyTorch tensors.
- **Model Definition:**
 - `torch.nn.Module` to define the MLP model.
 - `torch.nn.Linear` for fully connected layers.
 - `torch.nn.ModuleList` to manage layers dynamically.
 - Activation functions like `torch.nn.ReLU`, `torch.nn.Tanh`, etc.
- **Training Components:**
 - `torch.optim.Adam` for the optimizer.
 - `torch.nn.MSELoss` for the Mean Squared Error loss function.
- **Utilities:**
 - `torch.no_grad` to evaluate the model without computing gradients.

3 ****Bonus: Optuna Hyperparameter Tuning** (5 points)

You will implement hyperparameter tuning for the MLP Regressor using **Optuna**. The goal is to optimize key hyperparameters to achieve the best performance on the validation dataset and evaluate the tuned model on the test dataset.

1. **[2 Points]** Add a **Dropout** layer to your MLP Regressor architecture:
 - Incorporate `torch.nn.Dropout` to regularize the model and prevent overfitting.
 - The dropout rate should be one of the hyperparameters you tune.
2. **[3 Points]** Implement hyperparameter tuning using **Optuna**:
 - Optimize the following hyperparameters:
 - Learning rate.
 - Activation function (e.g., **ReLU**, **Tanh**).
 - Number of hidden layers.
 - Dropout rate.
 - Use the training and validation datasets to tune hyperparameters. Report every hyperparameter tuning trial.
 - Select the best hyperparameters based on the validation loss.
 - After finding the best hyperparameters, train the MLP Regressor using the training and validation datasets combined, and evaluate its performance on the test dataset. Report the following metrics:
 - Mean Squared Error (MSE)
 - Mean Absolute Error (MAE)